

## APM 598 Homework 1 - 6th February 2023

Siddharth Jain 1226137070

Ketan Choudhary 1226082301

Pranav Chougule 1225934595

```
In [ ]: # Required Imports
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]: # Dataset
filepath = 'data_HW1_ex1.csv'
data = np.loadtxt(filepath, delimiter=',')
```

```
In [ ]: # Splitting the dataset
x = data[:,0]
y = data[:,1]
k = 12 # given
```

```
In [ ]: # Loss function
def loss(degree, x, y):
    return np.sum((np.polyval(np.polyfit(x, y, degree), x) - y)**2)
```

```
In [ ]: def q1a():
    losses = []
    for degree in range(0, k):
        losses.append(loss(degree, x, y))
    plt.figure('Q 1a - Plot for Loss as a function of k')
    plt.plot(range(0, k), losses, 'r-')
    plt.xlabel('Degree')
    plt.ylabel('Loss')
    plt.show()
```

```
In [ ]: # Taking the first 80 data points for training and remaining 20 for test
x_train = data[0 : int(len(data) * 0.8), 0]
y_train = data[0 : int(len(data) * 0.8), 1]
x_test = data[int(len(data) * 0.8):, 0]
y_test = data[int(len(data) * 0.8):, 1]
```

```
In [ ]: def q1bc():
    error_train = []
    error_test = []

    for degree in range(0, k+1):
        error_train.append(loss(degree, x_train, y_train))
        poly_fit = np.poly1d(np.polyfit(x_train, y_train, degree))
        prediction = poly_fit(x_test)
        error_test.append(np.sum((prediction-y_test)**2))
    print('Degree   Train Loss   Test Loss')
    for degree, (er_train, er_test) in enumerate(zip(error_train, error_test)):
```

```

    print('%-9i%-13.4f%-4f'%(degree, er_train, er_test))
    k_star = error_test.index(min(error_test))
    print('\nOrder k* = {} since it has the least mean square error\n'.format(k_star))
    coefficient = np.polyfit(x, y, k_star)
    print('\nCoefficients for k*(', k_star, '): ', coefficient)

    fig = plt.figure('Q 1b - Test and Train Losses at Polynomial Degree k')
    diag1 = fig.add_subplot(1, 2, 1)
    training_line, = diag1.plot(range(0, k_star+1), error_train, c='k')
    test_line, = diag1.plot(range(0, k_star+1), error_test, c='r')
    diag1.legend([training_line, test_line], ["Train Loss", "Test Loss"])
    diag1.set(xlabel="Degree", ylabel="Loss")

    diag2 = fig.add_subplot(1, 2, 2)
    prediction = np.poly1d(np.polyfit(x, y, k_star))
    diag2.plot(x, y, 'ko')
    sort_x = sorted(x)
    fit_p, = diag2.plot(sort_x, prediction(sort_x), 'r-')
    diag2.set(xlabel = "x", ylabel = "y")
    diag2.legend([fit_p], ["K* = {}".format(k_star)])
    plt.show()


```

```

In [ ]: print('\nQ 1a\n')
        q1a()
        print('\nQ 1b 1c\n')
        q1bc()

```

Q 1a




Q 1b 1c

Degree	Train Loss	Test Loss
0	37.8110	36.0556
1	9.9070	2.3370
2	7.1658	5.8764
3	7.0453	8.5448
4	6.8797	15.9698
5	6.8533	21.7465
6	6.1257	22.6904
7	5.8454	26.9265
8	5.5408	240.0763
9	5.4494	3400.0992
10	5.3134	87689.9930
11	3.6145	18336610.0806
12	3.5227	57634507.3961

Order  $k^* = 1$  since it has the least mean square error

Coefficients for  $k^*(1)$ : [ 1.50632587 -1.35865816]



```
In [ ]: print('\nQ 2b 2c\n')
a0 = 1.8
b0 = 1.4
learning_rate = 0.05
gamma = 0.9
```

Q 2b 2c

```
In [ ]: def gradient_descent(x_train, y_train, a0, b0, learning_rate, numerically_optimal_a,
    iteration = 1
    n = len(x_train)
    old_theta = np.array([a0, b0])
    a = old_theta[0]
    b = old_theta[1]
    gradient_a = 1 / n * (np.sum(-2 * (y_train - a - (b * x_train))))
    gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - (a * x_train) - (b * (
    gradient_ab = np.array([gradient_a, gradient_b])
    new_theta = old_theta - (learning_rate * (gradient_ab))
    while((new_theta - old_theta).all() != 0):
        iteration += 1
        old_theta = new_theta
        a = old_theta[0]
        b = old_theta[1]
        gradient_a = 1 / n * (np.sum(-2 * (y_train - a - (b * x_train))))
        gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - (a * x_train) - (b * (
        gradient_ab = np.array([gradient_a, gradient_b])
        new_theta = old_theta - (learning_rate * gradient_ab)
    return [iteration, new_theta]
print('Method          Iterations    a_star    b_star')
iterations1, new_theta1 = gradient_descent(x_train, y_train, a0, b0, learning_rate,
print('Gradient Descent    %-14i%-10.4f%-4f'%(iterations1, new_theta1[0], new_theta1[1])
```

Method	Iterations	a_star	b_star
Gradient Descent	3459	-1.2421	1.4386

```
In [ ]: def momentum(x_train, y_train, gamma, a0, b0, learning_rate):
    iteration = 1
    n = len(x_train)
    old_theta = np.array([a0, b0])
    vn = np.array([0.0, 0.0])
    a, b = old_theta[0], old_theta[1]
    gradient_a = 1 / n * (np.sum(-2 * (y_train - a - (b * x_train))))
    gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - (a * x_train) - (b * (
    gradient_ab = np.array([gradient_a, gradient_b])
    new_vn = (gamma * vn) + (learning_rate * gradient_ab)
    new_theta = old_theta - new_vn
    while((new_theta - old_theta).all() != 0):
        iteration += 1
        old_theta = new_theta
        vn = new_vn
        a = old_theta[0]
        b = old_theta[1]
        gradient_a = 1 / n * (np.sum(-2 * (y_train - a - (b * x_train))))
        gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - (a * x_train) - (b * (
        gradient_ab = np.array([gradient_a, gradient_b])
        new_vn = (gamma * vn) + (learning_rate * gradient_ab)
        new_theta = old_theta - new_vn
    return [iteration, new_theta]
print('Method          Iterations    a_star    b_star')
iterations2, new_theta_momentum2 = momentum(x_train, y_train, gamma, a0, b0, learning_rate)
print('Momentum          %-14i%-10.4f%-4f'%(iterations2, new_theta_momentum2[0], new_theta_momentum2[1])
```

Method	Iterations	a_star	b_star
Momentum	593	-1.2421	1.4386

```
In [ ]: def nesterov(x_train, y_train, gamma, a0, b0, learning_rate):
    iteration = 1
    n = len(x_train)
    old_theta = np.array([a0, b0])
    a, b = old_theta[0], old_theta[1]
    vn = np.array([0.0, 0.0])
    gradient_a = 1 / n * (np.sum(-2 * (y_train - a - (gamma * vn[0])) - ((b - (gamma
    gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - ((a - (gamma * vn[0]))
    gradient_ab = np.array([gradient_a, gradient_b])
    new_vn = (gamma * vn) + (learning_rate * gradient_ab)
    new_theta = old_theta - new_vn
    while((new_theta - old_theta).all() != 0):
        iteration += 1
        old_theta = new_theta
        vn = new_vn
        a = old_theta[0]
        b = old_theta[1]
        gradient_a = 1 / n * (np.sum(-2 * (y_train - (a - (gamma * vn[0])) - ((b - (
        gradient_b = 1 / n * (np.sum(-2 * ((x_train * y_train) - ((a - (gamma * vn[
        gradient_ab = np.array([gradient_a, gradient_b])
        new_vn = (gamma * vn) + (learning_rate * gradient_ab)
        new_theta = old_theta - new_vn
    return [iteration, new_theta]
print('Method          Iterations    a_star    b_star')
iterations3, new_theta_nesterov3 = nesterov(x_train, y_train, gamma, a0, b0, learni
print('Nesterov          %-14i%-10.4f%-4f'%(iterations3, new_theta_nesterov3[0], n
```

Method	Iterations	a_star	b_star
Nesterov	553	-1.2421	1.4386

```
In [ ]: print("\nQ 2a\n")
x_train = data[0 : int(len(data)), 0]
y_train = data[0 : int(len(data)), 1]
n = len(x_train)
b_star = (np.sum(x_train * y_train) - (1 / n * (np.sum(x_train) * np.sum(y_train)))
a_star = 1 / n * (np.sum(y_train) - (b_star * np.sum(x_train)))
optimal_ab = np.array([a_star, b_star])
print('Deducing numerically:')
print('a_star: {}\nb_star: {}'.format(optimal_ab[0], optimal_ab[1]))
```

Q 2a

```
Deducing numerically:
a_star: -1.358658160753337
b_star: 1.5063258659470231
```

Q2a.

$$l(a, b) = \frac{1}{N} \sum_{i=1}^N |y_i - (a + bx_i)|^2$$

$$\therefore l(a, b) = \frac{1}{N} \sum_{i=1}^N |y_i - a - bx_i|^2$$

Partial derivative w.r.t a and b

$$\therefore \frac{\partial l}{\partial a} = \frac{1}{N} \sum_{i=1}^N [2(y_i - a - bx_i)(-1)] = 0$$

$$= -\frac{1}{N} \sum_{i=1}^N y_i + \frac{1}{N} \sum_{i=1}^N a + \frac{1}{N} \sum_{i=1}^N bx_i = 0$$

$$\therefore \frac{1}{N} \sum_{i=1}^N y_i = \frac{1}{N} \sum_{i=1}^N a + \frac{1}{N} \sum_{i=1}^N bx_i \quad \text{--- (I)}$$

$$\frac{\partial l}{\partial b} = \frac{1}{N} \sum_{i=1}^N [2(y_i - a - bx_i)(-x_i)] = 0$$

$$= -\frac{1}{N} \sum_{i=1}^N y_i x_i + \frac{1}{N} \sum_{i=1}^N ax_i + \frac{1}{N} \sum_{i=1}^N bx_i^2$$

$$\therefore \frac{1}{N} \sum_{i=1}^N y_i x_i = \frac{1}{N} \sum_{i=1}^N ax_i + \frac{1}{N} \sum_{i=1}^N bx_i^2 \quad \text{--- (II)}$$

$$\text{If } w = \frac{1}{N} \sum_{i=1}^N x_i \quad x = \frac{1}{N} \sum_{i=1}^N y_i$$

$$y = \frac{1}{N} \sum_{i=1}^N x_i y_i \quad z = \frac{1}{N} \sum_{i=1}^N x_i^2$$

$$\therefore y = a w + b z \quad \text{and} \quad x = a + b w \quad \text{--- (A)}$$

$$\text{Dividing (II) by } w \Rightarrow y/w = a + b z/w \quad \text{--- (III) (B)}$$

$$\therefore a = x - bw \quad \text{and} \quad b = \frac{y - xw}{z - ww}$$

$$\therefore a = \frac{1}{N} \sum_{i=1}^N y_i - \frac{1}{N} \sum_{i=1}^N x_i$$

$$b = \frac{\frac{1}{N} \sum_{i=1}^N x_i y_i - \left( \frac{1}{N} \sum_{i=1}^N x_i \right) \left( \frac{1}{N} \sum_{i=1}^N y_i \right)}{\frac{1}{N} \sum_{i=1}^N x_i^2 - \left( \frac{1}{N} \sum_{i=1}^N x_i \right) \left( \frac{1}{N} \sum_{i=1}^N x_i \right)}$$

```
In [ ]: import torch
        from torch import nn
        from torchvision import transforms
        from torchvision import datasets
        from torch.utils.data import DataLoader
        import numpy as np
        import matplotlib.pyplot as plt
```

```
In [ ]: cfg = dict()
        cfg['numEpochs'] = 40
        cfg['learning_rate'] = 0.0001
        cfg['batchSize'] = 4
```

```
In [ ]: MNIST_training_set_tensor = datasets.FashionMNIST('data', train = True, download =
        MNIST_test_set_tensor = datasets.FashionMNIST('data', train = False, download = True)
        myLoader_train = DataLoader(MNIST_training_set_tensor, shuffle = True, batch_size =
        myLoader_test = DataLoader(MNIST_test_set_tensor, shuffle = False, batch_size = cfg
```

```
In [ ]: #myModel
        myModel = nn.Linear(784, 10)
        optimizer = torch.optim.Adam(myModel.parameters(), lr = cfg['learning_rate'])
        myLoss = nn.CrossEntropyLoss()
        #print(myModel.state_dict())
        final_loss=[]
```

```
In [ ]: print("Q 3a\n")
        for epoch in range(cfg['numEpochs']):
            running_loss = 0
            minibatch = 0
            for x_mini, y_mini in myLoader_train:
                optimizer.zero_grad()
                N, _, q, r = x_mini.size()
                score = myModel(x_mini.view(N,q*r))
                loss = myLoss(score,y_mini)
                loss.backward()
                optimizer.step()
                running_loss += loss.detach().numpy()
                minibatch += 1
            average_loss = running_loss/minibatch
```



```
final_loss.append(average_loss)
print("E:{} Average Loss: {}".format(epoch+1, average_loss))
```


Q 3a

```
E:1 Average Loss: 0.7139170745755236
E:2 Average Loss: 0.5159314553940514
E:3 Average Loss: 0.4774667513840599
E:4 Average Loss: 0.45732658601621323
E:5 Average Loss: 0.44381647969970167
E:6 Average Loss: 0.43416462090292673
E:7 Average Loss: 0.42735239889814014
E:8 Average Loss: 0.42137811187998936
E:9 Average Loss: 0.4171578940883822
E:10 Average Loss: 0.41254373762860974
E:11 Average Loss: 0.40950242257257147
E:12 Average Loss: 0.40614574913239143
E:13 Average Loss: 0.4036863504893709
E:14 Average Loss: 0.4013412846794924
E:15 Average Loss: 0.3994920930537085
E:16 Average Loss: 0.3975230125461181
E:17 Average Loss: 0.3955777901786894
E:18 Average Loss: 0.3936058884898977
E:19 Average Loss: 0.3920777539262936
E:20 Average Loss: 0.3908284713284559
E:21 Average Loss: 0.38955447497993884
E:22 Average Loss: 0.3881976457363664
E:23 Average Loss: 0.38711647709034924
E:24 Average Loss: 0.3861849814362533
E:25 Average Loss: 0.38527873083345865
E:26 Average Loss: 0.38439105056159123
E:27 Average Loss: 0.38300268993816766
E:28 Average Loss: 0.38275217127406747
E:29 Average Loss: 0.3817324692709752
E:30 Average Loss: 0.38088271303217236
E:31 Average Loss: 0.3799307055525336
E:32 Average Loss: 0.37938072058312344
E:33 Average Loss: 0.378750729895388
E:34 Average Loss: 0.37795806020691913
E:35 Average Loss: 0.3776280917688588
E:36 Average Loss: 0.3770422859154244
E:37 Average Loss: 0.3762354942224425
E:38 Average Loss: 0.3760144658959796
E:39 Average Loss: 0.37513915927561997
E:40 Average Loss: 0.3745972739414162
```

```
In [ ]: print('\nQ 3b\n')
weights = myModel.state_dict()['weight']
plt.plot([i for i in range(cfg['numEpochs']) ],final_loss)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.show()
plt.clf()
```

Q 3b





<Figure size 640x480 with 0 Axes>

```
In [ ]: print("\nQ 3c\n")
label_fashion = ['T-shirt', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shir
for i in range(10):
    plt.clf()
    plt.imshow(weights[i].view(28,28), vmin = -0.5, vmax = 0.5, cmap = 'seismic')
    plt.title('After Training, Template for {}'.format(label_fashion[i]))
    plt.colorbar(extend='both')
    plt.show()
```

Q 3c

